

The background of the image is a dense, close-up photograph of variegated plant leaves, likely a Philodendron. The leaves are a deep green color with prominent, lighter green or yellowish-white veins and edges, creating a complex, textured pattern. The lighting is soft, highlighting the natural patterns of the foliage.

Plant species classification

Svanidze Mariam

My data

47 plant species:

Calathea

Kalanchoe

Ficus elastica

Money Tree

Aloe Vera

Dracaena

Schefflera

Orchid

Yucca

Tulip

etc





Data processing

Split data to train and test sets taking 80% and 20% corresponding percentage from each class

```
split = StratifiedShuffleSplit(n_splits=1, test_size=0.2, random_state=42)
```

```
train_val_idx, test_idx = next(split.split(np.zeros(len(labels)), labels))
```

UserWarning: Palette images with Transparency expressed in bytes should be converted to RGBA images

Solution:

```
convert_palette_images_to_rgb('/content/house_plant_species/house_plant_species', save_new=False)
```

```
delete_corrupted_images('/content/house_plant_species/house_plant_species')
```

Add Kfold Cross-validation:

```
for fold, (train_idx, val_idx) in enumerate(skf.split(np.zeros(len(labels_train_val)), labels_train_val)):
```

```
==== Fold 1 ==== Fold 2 ==== Fold 3 ==== Fold 4 ==== Fold 5 ===
```

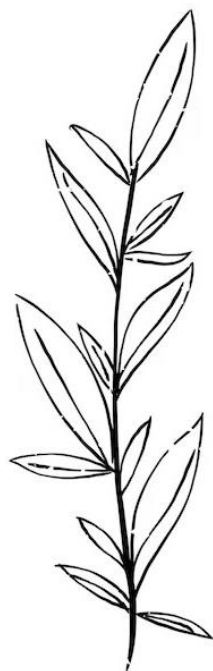
Correct test and train sets:

```
for fold, (train_idx, val_idx) in enumerate(skf.split(np.zeros(len(labels_train_val)), labels_train_val)):
```

```
    train_subset = Subset(train_val_subset, train_idx)
```

```
    val_subset = Subset(train_val_subset, val_idx)
```

Final Result



```
Train/Val classes: [ 0  1  2  3  4  5  6  7  8  9 10 11 12 13 14 15 16 17 18 19 20 21 22 23
                   24 25 26 27 28 29 30 31 32 33 34 35 36 37 38 39 40 41 42 43 44 45 46]
Test classes:      [ 0  1  2  3  4  5  6  7  8  9 10 11 12 13 14 15 16 17 18 19 20 21 22 23
                   24 25 26 27 28 29 30 31 32 33 34 35 36 37 38 39 40 41 42 43 44 45 46]
```

=== Fold 1 ===

```
Train classes: [ 0  1  2  3  4  5  6  7  8  9 10 11 12 13 14 15 16 17 18 19 20 21 22 23
                24 25 26 27 28 29 30 31 32 33 34 35 36 37 38 39 40 41 42 43 44 45 46]
Val classes:   [ 0  1  2  3  4  5  6  7  8  9 10 11 12 13 14 15 16 17 18 19 20 21 22 23
                24 25 26 27 28 29 30 31 32 33 34 35 36 37 38 39 40 41 42 43 44 45 46]
```

=== Fold 2 ===

```
Train classes: [ 0  1  2  3  4  5  6  7  8  9 10 11 12 13 14 15 16 17 18 19 20 21 22 23
                24 25 26 27 28 29 30 31 32 33 34 35 36 37 38 39 40 41 42 43 44 45 46]
Val classes:   [ 0  1  2  3  4  5  6  7  8  9 10 11 12 13 14 15 16 17 18 19 20 21 22 23
                24 25 26 27 28 29 30 31 32 33 34 35 36 37 38 39 40 41 42 43 44 45 46]
```

=== Fold 3 ===

```
Train classes: [ 0  1  2  3  4  5  6  7  8  9 10 11 12 13 14 15 16 17 18 19 20 21 22 23
                24 25 26 27 28 29 30 31 32 33 34 35 36 37 38 39 40 41 42 43 44 45 46]
Val classes:   [ 0  1  2  3  4  5  6  7  8  9 10 11 12 13 14 15 16 17 18 19 20 21 22 23
                24 25 26 27 28 29 30 31 32 33 34 35 36 37 38 39 40 41 42 43 44 45 46]
```

=== Fold 4 ===

```
Train classes: [ 0  1  2  3  4  5  6  7  8  9 10 11 12 13 14 15 16 17 18 19 20 21 22 23
                24 25 26 27 28 29 30 31 32 33 34 35 36 37 38 39 40 41 42 43 44 45 46]
Val classes:   [ 0  1  2  3  4  5  6  7  8  9 10 11 12 13 14 15 16 17 18 19 20 21 22 23
                24 25 26 27 28 29 30 31 32 33 34 35 36 37 38 39 40 41 42 43 44 45 46]
```

=== Fold 5 ===

```
Train classes: [ 0  1  2  3  4  5  6  7  8  9 10 11 12 13 14 15 16 17 18 19 20 21 22 23
                24 25 26 27 28 29 30 31 32 33 34 35 36 37 38 39 40 41 42 43 44 45 46]
Val classes:   [ 0  1  2  3  4  5  6  7  8  9 10 11 12 13 14 15 16 17 18 19 20 21 22 23
                24 25 26 27 28 29 30 31 32 33 34 35 36 37 38 39 40 41 42 43 44 45 46]
```

My first model

Epoch [1/10], Loss: 3.3205, Train Acc: 12.29%, Test Acc: 20.07%

Epoch [2/10], Loss: 2.7619, Train Acc: 24.66%, Test Acc: 26.23%

Epoch [3/10], Loss: 2.3523, Train Acc: 34.03%, Test Acc: 29.54%

Epoch [4/10], Loss: 1.9357, Train Acc: 45.34%, Test Acc: 31.13%

Epoch [5/10], Loss: 1.5318, Train Acc: 55.89%, Test Acc: 30.93%

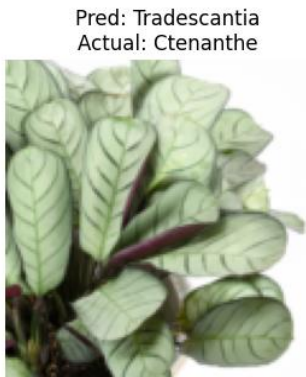
Epoch [6/10], Loss: 1.1150, Train Acc: 67.15%, Test Acc: 29.48%

Epoch [7/10], Loss: 0.7394, Train Acc: 78.32%, Test Acc: 29.54%

Epoch [8/10], Loss: 0.4189, Train Acc: 87.69%, Test Acc: 28.60%

Epoch [9/10], Loss: 0.2441, Train Acc: 92.76%, Test Acc: 28.70%

Epoch [10/10], Loss: 0.1371, Train Acc: 96.15%, Test Acc: 27.72%



Pred: Snake plant (Sansevieria)
Actual: Jade plant (Crassula ovata)



```
def MScnn(num_classes):  
    # Feature extraction  
    feature_extractor = nn.Sequential(  
        nn.Conv2d(3, 6, kernel_size=5),      # Output: (128 - 5 + 1) = 124 → 6x124x124  
        nn.ReLU(),  
        nn.MaxPool2d(2, 2),                  # → 6x62x62  
        nn.Conv2d(6, 16, kernel_size=5),     # → (62 - 5 + 1) = 58 → 16x58x58  
        nn.ReLU(),  
        nn.MaxPool2d(2, 2)                   # → 16x29x29  
    )  
  
    # Fully connected classifier  
    classifier = nn.Sequential(  
        nn.Linear(16 * 29 * 29, 120),         # Flattened features  
        nn.ReLU(),  
        nn.Linear(120, 84),  
        nn.ReLU(),  
        nn.Linear(84, num_classes)           # Output layer  
    )  
  
    # Forward pass function  
    def forward_function(x):  
        x = feature_extractor(x)  
        x = torch.flatten(x, 1) # flatten all except batch  
        x = classifier(x)  
        return x  
  
    # Build full model  
    model = nn.Sequential(feature_extractor, nn.Flatten(), classifier)  
    model.forward = forward_function  
  
    return model
```

My second model



```
class ThirdMScnn(nn.Module):
    def __init__(self, num_classes):
        super(ThirdMScnn, self).__init__()

        self.feature_extractor = nn.Sequential(
            nn.Conv2d(3, 6, kernel_size=5),      # → 6x124x124
            nn.BatchNorm2d(6),
            nn.ReLU(),
            nn.MaxPool2d(2, 2),                  # → 6x62x62
            nn.Dropout(0.25),

            nn.Conv2d(6, 16, kernel_size=3),     # → 16x60x60
            nn.BatchNorm2d(16),
            nn.ReLU(),
            nn.MaxPool2d(2, 2),                  # → 16x30x30
            nn.Dropout(0.25),
        )

        self.classifier = nn.Sequential(
            nn.Linear(16 * 30 * 30, 120),
            nn.ReLU(),

            nn.Linear(120, 84),
            nn.ReLU(),

            nn.Linear(84, num_classes)
        )
```

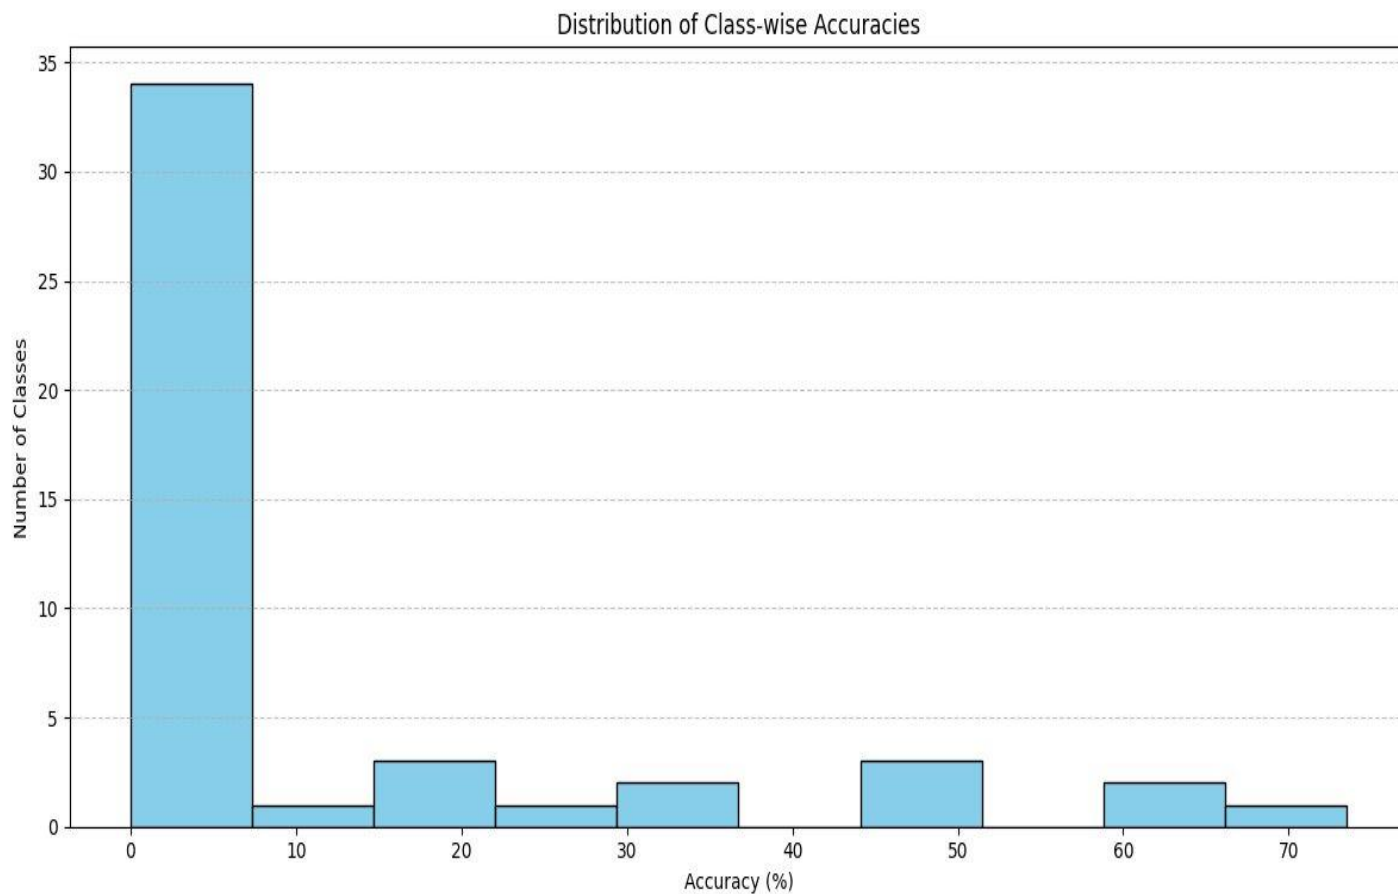
Adding data augmentation

```
train_transform = transforms.Compose([
    transforms.Resize((128, 128), interpolation=InterpolationMode.BILINEAR),
    transforms.RandomAffine(degrees=45, translate=(0.1, 0.1), shear=15),
    transforms.RandomHorizontalFlip(),
    transforms.ColorJitter(brightness=0.9, contrast=0.9),
    transforms.ToTensor(),
    transforms.Normalize(mean, std)
])
```

Weight initialization

```
def init_weights(m):
    if isinstance(m, nn.Conv2d):
        nn.init.kaiming_normal_(m.weight, mode='fan_out', nonlinearity='relu')
        if m.bias is not None:
            nn.init.constant_(m.bias, 0)
    elif isinstance(m, nn.Linear):
        nn.init.kaiming_normal_(m.weight)
        nn.init.constant_(m.bias, 0)
    elif isinstance(m, nn.BatchNorm2d):
        nn.init.constant_(m.weight, 1)
        nn.init.constant_(m.bias, 0)
```


Class accuracy distribution

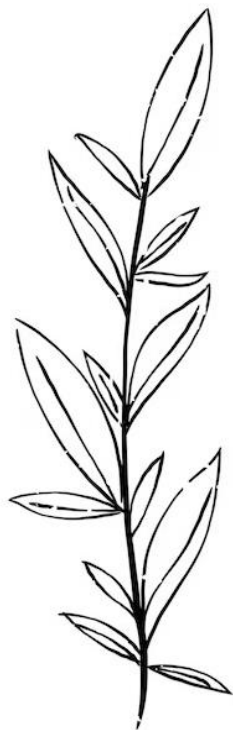


Epoch [1/20], Loss: 3.7200, Train Acc: 3.81%, Val Acc: 6.56%
Epoch [2/20], Loss: 3.4464, Train Acc: 6.82%, Val Acc: 8.51%
Epoch [3/20], Loss: 3.4120, Train Acc: 8.54%, Val Acc: 9.44%
Epoch [4/20], Loss: 3.1666, Train Acc: 9.25%, Val Acc: 12.06%
Epoch [5/20], Loss: 3.1588, Train Acc: 9.62%, Val Acc: 11.30%
Epoch [6/20], Loss: 3.0785, Train Acc: 10.69%, Val Acc: 12.23%
Epoch [7/20], Loss: 3.1652, Train Acc: 11.03%, Val Acc: 12.95%
Epoch [8/20], Loss: 3.1107, Train Acc: 11.55%, Val Acc: 12.40%
Epoch [9/20], Loss: 3.1090, Train Acc: 11.34%, Val Acc: 13.12%
Epoch [10/20], Loss: 3.0969, Train Acc: 12.01%, Val Acc: 13.80%
Final Test Accuracy: 14.92% for 20 epochs



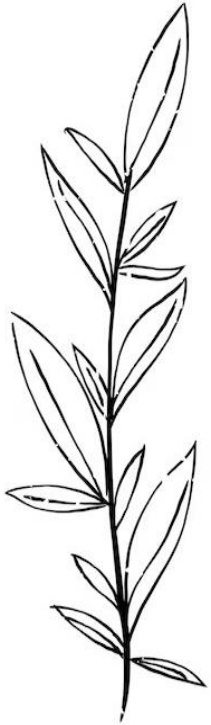
Zero accuracy model

Epoch 1/10 – Train Loss: 3.2668 – Test Accuracy: 21.52%
Epoch 2/10 – Train Loss: 2.6699 – Test Accuracy: 25.18%
Epoch 3/10 – Train Loss: 2.2177 – Test Accuracy: 30.63%
Epoch 4/10 – Train Loss: 1.7280 – Test Accuracy: 32.39%
Epoch 5/10 – Train Loss: 1.2307 – Test Accuracy: 33.10%
Epoch 6/10 – Train Loss: 0.7755 – Test Accuracy: 31.24%
Epoch 7/10 – Train Loss: 0.4228 – Test Accuracy: 30.59%
Epoch 8/10 – Train Loss: 0.2082 – Test Accuracy: 30.49%
Epoch 9/10 – Train Loss: 0.1167 – Test Accuracy: 29.00%
Epoch 10/10 – Train Loss: 0.0786 – Test Accuracy: 29.54%



Accuracy for class: African Violet (*Saintpaulia ionantha*) is 28.4 %
Accuracy for class: Aloe Vera is 0.0 %
Accuracy for class: Anthurium (*Anthurium andraeanum*) is 1.1 %
Accuracy for class: Areca Palm (*Dypsis lutescens*) is 0.0 %
Accuracy for class: Asparagus Fern (*Asparagus setaceus*) is 0.0 %
Accuracy for class: Begonia (*Begonia* spp.) is 0.0 %
Accuracy for class: Bird of Paradise (*Strelitzia reginae*) is 0.0 %
Accuracy for class: Birds Nest Fern (*Asplenium nidus*) is 0.0 %
Accuracy for class: Boston Fern (*Nephrolepis exaltata*) is 0.0 %
Accuracy for class: Calathea is 1.5 %
Accuracy for class: Cast Iron Plant (*Aspidistra elatior*) is 0.0 %
Accuracy for class: Chinese Money Plant (*Pilea peperomioides*) is 0.0 %
Accuracy for class: Chinese evergreen (*Aglaonema*) is 7.8 %
Accuracy for class: Christmas Cactus (*Schlumbergera bridgesii*) is 1.6 %
Accuracy for class: Chrysanthemum is 0.0 %
Accuracy for class: Ctenanthe is 2.9 %
Accuracy for class: Daffodils (*Narcissus* spp.) is 45.2 %
Accuracy for class: Dracaena is 0.0 %
Accuracy for class: Dumb Cane (*Dieffenbachia* spp.) is 62.0 %
Accuracy for class: Elephant Ear (*Alocasia* spp.) is 0.0 %
Accuracy for class: English Ivy (*Hedera helix*) is 0.0 %
Accuracy for class: Hyacinth (*Hyacinthus orientalis*) is 32.8 %
Accuracy for class: Iron Cross begonia (*Begonia masoniana*) is 1.9 %
Accuracy for class: Jade plant (*Crassula ovata*) is 18.3 %
Accuracy for class: Kalanchoe is 0.0 %
Accuracy for class: Lilium (*Hemerocallis*) is 65.6 %
Accuracy for class: Lily of the valley (*Convallaria majalis*) is 73.5 %
Accuracy for class: Money Tree (*Pachira aquatica*) is 0.0 %
Accuracy for class: Monstera Deliciosa (*Monstera deliciosa*) is 20.0 %
Accuracy for class: Orchid is 6.4 %
Accuracy for class: Parlor Palm (*Chamaedorea elegans*) is 0.0 %
Accuracy for class: Peace lily is 6.5 %
Accuracy for class: Poinsettia (*Euphorbia pulcherrima*) is 50.8 %
Accuracy for class: Polka Dot Plant (*Hypoestes phylllostachya*) is 48.5 %
Accuracy for class: Ponytail Palm (*Beaucarnea recurvata*) is 0.0 %
Accuracy for class: Pothos (*Ivy arum*) is 0.0 %
Accuracy for class: Prayer Plant (*Maranta leuconeura*) is 0.0 %
Accuracy for class: Rattlesnake Plant (*Calathea lancifolia*) is 3.2 %
Accuracy for class: Rubber Plant (*Ficus elastica*) is 0.0 %
Accuracy for class: Sago Palm (*Cycas revoluta*) is 0.0 %
Accuracy for class: Schefflera is 0.0 %
Accuracy for class: Snake plant (*Sansevieria*) is 6.3 %
Accuracy for class: Tradescantia is 33.8 %
Accuracy for class: Tulip is 0.0 %
Accuracy for class: Venus Flytrap is 5.0 %
Accuracy for class: Yucca is 0.0 %
Accuracy for class: ZZ Plant (*Zamioculcas zamiifolia*) is 21.6 %

My third model



Selected classes:

Begonia

Calathea

Dracaena

Sansevieria

Ctenanthe

Aglaonema

Anthurium

Tradescantia

Monstera Deliciosa

Maranta leuconeura

Dumb Cane (Dieffenbachia)

Money Tree (Pachira aquatica)

```
def MScnn(num_classes):
    feature_extractor = nn.Sequential(
        nn.Conv2d(3, 6, kernel_size=5),      # → 220x220
        nn.ReLU(),
        nn.MaxPool2d(2, 2),                  # → 110x110
        nn.Conv2d(6, 16, kernel_size=5),     # → 106x106
        nn.ReLU(),
        nn.MaxPool2d(2, 2)                   # → 53x53
    )

    classifier = nn.Sequential(
        nn.Linear(16 * 53 * 53, 120),
        nn.ReLU(),
        nn.Linear(120, 84),
        nn.ReLU(),
        nn.Linear(84, num_classes)
    )

    def forward_function(x):
        x = feature_extractor(x)
        x = torch.flatten(x, 1)
        x = classifier(x)
        return x

    model = nn.Sequential(feature_extractor, nn.Flatten(), classifier)
    model.forward = forward_function
    return model
```

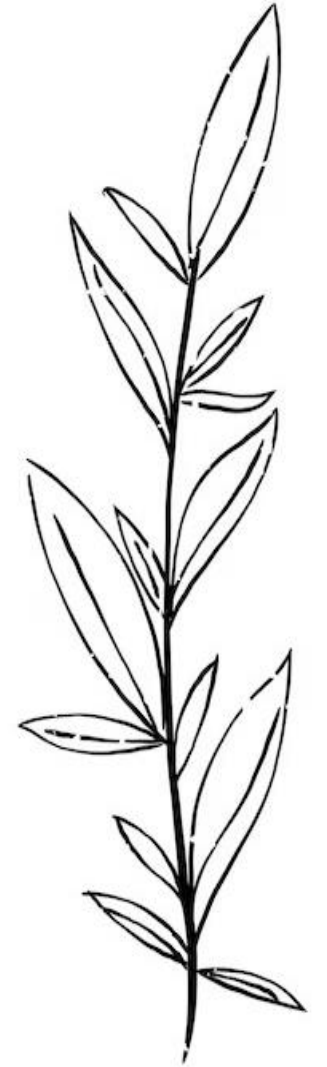
My third model

Selected classes:

Begonia
Calathea
Dracaena
Sansevieria
Ctenanthe
Aglaonema
Anthurium
Tradescantia
Monstera Deliciosa
Maranta leuconeura
Dumb Cane (Dieffenbachia)
Money Tree (Pachira aquatica)

Outcome:

Epoch [1/10], Loss: 2.4640, Train Acc: 10.75%, Test Acc: 18.86%
Epoch [2/10], Loss: 2.2097, Train Acc: 23.09%, Test Acc: 25.85%
Epoch [3/10], Loss: 2.0342, Train Acc: 28.73%, Test Acc: 27.97%
Epoch [4/10], Loss: 1.9165, Train Acc: 33.02%, Test Acc: 29.87%
Epoch [5/10], Loss: 1.8056, Train Acc: 36.68%, Test Acc: 29.66%
Epoch [6/10], Loss: 1.6470, Train Acc: 42.03%, Test Acc: 32.42%
Epoch [7/10], Loss: 1.4319, Train Acc: 50.08%, Test Acc: 32.94%
Epoch [8/10], Loss: 1.0327, Train Acc: 64.22%, Test Acc: 32.63%
Epoch [9/10], Loss: 0.6522, Train Acc: 77.09%, Test Acc: 33.47%
Epoch [10/10], Loss: 0.3434, Train Acc: 88.08%, Test Acc: 33.47%



Results

Train Accuracy: 94.57%

Test Accuracy: 33.47%

Accuracy per class

Begonia : 97.85%

Begonia: 19.15%

Calathea : 97.35%

Calathea : 18.18%

Dracaena : 100.00%

Dracaena : 26.92%

Sansevieria : 97.16%

Sansevieria: 46.84%

Ctenanthe : 98.18%

Ctenanthe : 28.99%

Aglaonema: 85.64%

Aglaonema : 14.56%

Anthurium : 96.69%

Anthurium: 60.44%

Tradescantia : 97.80%

Tradescantia : 50.00%

Monstera Deliciosa : 94.75%

Monstera Deliciosa: 34.86%

Maranta leuconeura : 95.62%

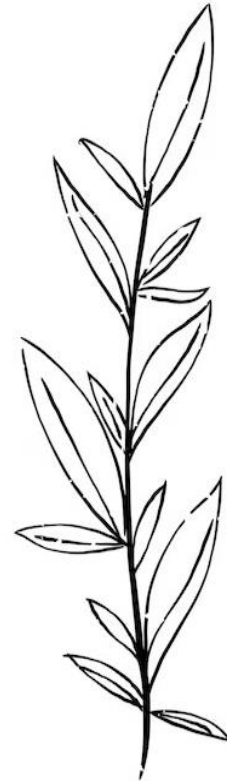
Maranta leuconeura: 21.25%

Dumb Cane (Dieffenbachia) : 85.45%

DumbCane (Dieffenbachia) : 28.70%

Money Tree (Pachira aquatica) : 98.95%

Money Tree (Pachira aquatica) : 47.22%



Pre-trained model example

```
# Load pretrained VGG16 and modify for 47 classes
vgg16 = models.vgg16(pretrained=True)
for param in vgg16.parameters():
    param.requires_grad = False # Freeze feature extractor

# Replace final layer
vgg16.classifier[6] = nn.Linear(4096, 47)
vgg16 = vgg16.to(device)

# Loss and optimizer
criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(vgg16.classifier[6].parameters(), lr=0.001)
```

Results:

Epoch [1/10],	Loss: 408.5554,	Train Acc: 27.68%
Epoch [2/10],	Loss: 334.3769,	Train Acc: 37.95%
Epoch [3/10],	Loss: 326.3154,	Train Acc: 40.38%
Epoch [4/10],	Loss: 318.5549,	Train Acc: 41.23%



Data issues



Unrealistic images



Money tree (*Pachira aquatica*)



Unbalanced classes

Prayer Plant
(Mara...eura)
400 items
Shared by Me

Rattlesnake
Plant...ifolia)
316 items
Shared by Me

Rubber Plant
(Ficus...stica)
291 items
Shared by Me

Sago Palm
(Cyca...oluta)
202 items
Shared by Me

Schefflera
326 items
Shared by Me

Tradescantia
341 items
Shared by Me

Tulip
341 items
Shared by Me

Venus Flytrap
199 items
Shared by Me

Yucca
66 items
Shared by Me

ZZ Plant
(Zami...iifolia)
438 items
Shared by Me

African Violet
(Saint...nthia)
337 items
Shared by Me

Aloe Vera
252 items
Shared by Me

Anthurium
(Anth...num)
455 items
Shared by Me

Areca Palm
(Dyps...cens)
189 items
Shared by Me

Asparagus Fern
(Aspa...ceus)
169 items
Shared by Me

Begonia
(Bego...spp.)
236 items
Shared by Me

Bird of Paradise
(Strel...inae)
180 items
Shared by Me

Birds Nest Fern
(Aspl...idus)
290 items
Shared by Me

Boston Fern
(Neph...ltata)
307 items
Shared by Me

Calathea
330 items
Shared by Me

Cast Iron Plant
(Aspi...lating)
266 items
Shared by Me

Chinese
ever...ema)
514 items
Shared by Me

Chinese Money
Plant...ides)
382 items
Shared by Me

Christmas
Cactu...gesii)
312 items
Shared by Me

Chrysanthem
um
209 items
Shared by Me

Ctenanthe
347 items
Shared by Me

Daffodils
(Narci...spp.)
421 items
Shared by Me

Dracaena
261 items
Shared by Me

252. 455. 189. 169. 180. 307. 514. 209. 421. 261.

400. 316. 291. 202. 326. 341. 341. 199. 66. 438.

Begonia



Variegation within a plant species

Antherium





Similarities across different species

Marantaceae family

Maranta



Calathea



Ctenanthe





Every seed planted today is tomorrow's harvest.